

BANK LOAN APPROVAL PREDICTION - LOGISTICS REGRESSION PROJECT

Predict whether a bank should approve a customer's loan Goal: Predict Approved = 1 or Not Approved = 0 using customer features.

Project Covers:

Clean dataset loading

Feature selection

Train/test split

Scaling

Training(logistic regression)

Full evaluation (Accuracy, Precision, Recall, F1)

Confusion Matrix + Classification Report

Decision thresholds

AUC - ROC curve

Prediction function

Professional project structure

```
# pandas: Used for data handling and analysis
import pandas as pd

# numpy: Used for numerical operations
import numpy as np

# matplotlib.pyplot: Used for basic data visualization
import matplotlib.pyplot as plt

# seaborn: Built on matplotlib, used for advanced & attractive plots
import seaborn as sns

# train_test_split: Used to divide data into training & testing sets
# Ensures model is tested on unseen data
from sklearn.model_selection import train_test_split

# StandardScaler: Used to scale features (mean = 0, std = 1)
# Important for models like Logistic Regression
from sklearn.preprocessing import StandardScaler
```

```

# LogisticRegression: A classification algorithm
# Used to predict binary outcomes (Yes/No, 0/1)
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import (
    accuracy_score,      # Accuracy: Out of all predictions, how many
    were correct?
    precision_score,     # Precision: Out of all YES predictions, how
    many were actually YES?
    recall_score,        # Recall: Out of all actual YES cases, how
    many did the model find?
    f1_score,            # F1 Score: A single score that balances
Precision and Recall
    confusion_matrix,   # A table that shows correct and wrong
predictions (YES/NO)
    classification_report, # Shows Accuracy, Precision, Recall, and F1
together in one report
    roc_curve,          # A graph that shows how well the model
separates YES and NO
    roc_auc_score       # A number that tells how good the model is
at separating classes
)

# Load Dataset
df = pd.read_csv("bank_loan.csv")
print("Dataset Loaded Successfully")
print(df.head())
print("\nShape:", df.shape)

Dataset Loaded Successfully
   Age  Income  CreditScore EmploymentStatus  Approved
0   59    40358           812      Unemployed      0
1   49    23267           595      Unemployed      0
2   35    102745          619      Salaried       1
3   63    109588          871  Self-employed      0
4   28    58513           648      Salaried       1

Shape: (500, 5)

print("=====Data Overview
=====")
#Data types and info
print("\nDataset Info:")
df.info()
# Check missing values
print("\nMissing Values in Each Column:")
print(df.isnull().sum());

```

```

=====Data Overview=====
=====

Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              500 non-null    int64  
 1   Income            500 non-null    int64  
 2   CreditScore       500 non-null    int64  
 3   EmploymentStatus 500 non-null    object  
 4   Approved          500 non-null    int64  
dtypes: int64(4), object(1)
memory usage: 19.7+ KB

Missing Values in Each Column:
Age          0
Income        0
CreditScore   0
EmploymentStatus 0
Approved      0
dtype: int64

# Count each Employment Status Category

print("\nEmployment Status Distribution:")
print(df["EmploymentStatus"].value_counts())

# Preview the last 10 rows
print("\nLast 10 Rows of the Dataset:")
print(df.tail(10));

Employment Status Distribution:
EmploymentStatus
Salaried      289
Self-employed 136
Unemployed    75
Name: count, dtype: int64

Last 10 Rows of the Dataset:
   Age  Income  CreditScore EmploymentStatus  Approved
490  43     32219          870  Self-employed      0
491  37     20235          662  Self-employed      0
492  46     62929          856  Self-employed      0
493  28     147309         644   Salaried        1
494  49     87444          469   Salaried        1
495  46     159639         655   Salaried        1

```

496	30	121834	490	Salaried	0
497	46	104555	721	Salaried	1
498	54	93698	306	Unemployed	0
499	61	144450	432	Salaried	0

```
# Summary Statistics for numerical features
print("\nSummary Statistics:")
print(df.describe());

# Unique values in Employment Status
print("\nUnique Employment Status Values:")
print(df["EmploymentStatus"].unique());
```


Summary Statistics:

	Age	Income	CreditScore	Approved
count	500.000000	500.000000	500.000000	500.000000
mean	43.116000	109045.696000	583.490000	0.502000
std	12.733217	52974.205023	171.614805	0.500497
min	21.000000	20235.000000	300.000000	0.000000
25%	32.000000	62479.250000	440.000000	0.000000
50%	44.000000	109228.500000	585.500000	1.000000
75%	53.000000	156195.000000	723.500000	1.000000
max	64.000000	199208.000000	898.000000	1.000000

Unique Employment Status Values:

```
['Unemployed' 'Salaried' 'Self-employed']
```

```
# 4 . Employment Status Count Plot
plt.figure(figsize=(6,4))
sns.countplot(x="EmploymentStatus", data=df)
plt.title("Employment Status Distribution")
plt.xlabel("Employment Status")
plt.ylabel("Count")
plt.show()
```



FEATURE SELECTION & CLEANING

```
# Select important features
df = df[["Age", "Income", "CreditScore", "EmploymentStatus",
"Approved"]]
```

Encode Employment Status (numeric encoding)

```
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder();
df["EmploymentEncoded"] = le.fit_transform(df["EmploymentStatus"])

# Final Feature / Target
X = df[["Age", "Income", "CreditScore", "EmploymentEncoded"]]
y = df["Approved"];

print("\nFinal Features:\n", X.head());

Final Features:
    Age  Income  CreditScore  EmploymentEncoded
```

0	59	40358	812	2
1	49	23267	595	2
2	35	102745	619	0
3	63	109588	871	1
4	28	58513	648	0

TRAIN–TEST SPLIT

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42, stratify=y
    # stratify=y ensures the train and test sets have the same
    # proportion of 0 and 1, avoiding imbalance.
);
print("Train size:", X_train.shape);
print("Test size:", X_test.shape);

Train size: (375, 4)
Test size: (125, 4)
```

FEATURE SCALING

```
scaler = StandardScaler()
# Create a scaler that standardizes data (mean = 0, std = 1)

X_train_scaled = scaler.fit_transform(X_train)
# Learn scaling from training data and apply it

X_test_scaled = scaler.transform(X_test)
# Apply the SAME scaling to test data
```

TRAIN LOGISTIC REGRESSION MODEL

```
model = LogisticRegression(max_iter=1000) # Improves accuracy
# Create a Logistic Regression model

model.fit(X_train_scaled, y_train)
# Train the model using the scaled training data and their correct
# labels

print("Model Training Completed!")
# Confirm that training is finished

Model Training Completed!
```

MODEL PREDICTIONS

```
y_pred = model.predict(X_test_scaled)
# Predict class labels (0 or 1) for the test data

y_prob = model.predict_proba(X_test_scaled)[:, 1]
# Get probability of class 1 (loan approved) for each test sample
```

EVALUATION METRICS

```
print("\n===== MODEL EVALUATION =====")

print("Accuracy : ", accuracy_score(y_test, y_pred));
# How many total predictions were correct

print("Precision:", precision_score(y_test, y_pred));
# Out of all predicted YES, how many were actually YES (controls false positives)

print("Recall    : ", recall_score(y_test, y_pred));
# Out of all actual YES, how many the model correctly found (controls false negatives)

print("F1 Score : ", f1_score(y_test, y_pred));
# Balance of Precision and Recall in one score

===== MODEL EVALUATION =====
Accuracy : 0.624
Precision: 0.6212121212121212
Recall    : 0.6507936507936508
F1 Score : 0.6356589147286822
```

Model Evaluation

```
Accuracy = 0.62 (62%)
Out of 100 loan predictions, about 62 are correct.

Precision = 0.62 (62%)
When the model says "YES – loan approved", it is correct 62 times out of 100.

Recall = 0.65 (65%)
Out of all people who should get the loan, the model correctly finds 65% of them.
```

```
F1 Score = 0.63 (63%)
Overall balanced performance between Precision and Recall.
```

```
Percentages help us understand model performance easily, regardless of
dataset size.
```

CONFUSION MATRIX + CLASSIFICATION REPORT

```
print("\n===== CONFUSION MATRIX =====")
print(confusion_matrix(y_test, y_pred));
# Shows TP, FP, TN, FN to understand correct and incorrect predictions

===== CONFUSION MATRIX =====
[[37 25]
 [22 41]]
```

Confusion Matrix

```
37 → Correctly predicted NO
25 → Wrongly predicted YES (but actually NO)
22 → Wrongly predicted NO (but actually YES)
41 → Correctly predicted YES
```

```
print("\n===== CLASSIFICATION REPORT =====");
print(classification_report(y_test, y_pred));
# Shows Precision, Recall, F1-score, and support for each class
```

```
===== CLASSIFICATION REPORT =====
      precision    recall   f1-score   support
          0         0.63     0.60     0.61       62
          1         0.62     0.65     0.64       63

   accuracy                           0.62      125
  macro avg       0.62     0.62     0.62      125
weighted avg     0.62     0.62     0.62      125
```

Class 0 (NO / Loan Rejected)

Precision 63% → When model says NO, it is correct 63 times out of 100
Recall 60% → Found 60% of actual NO cases
F1-score 61% → Balanced score for NO class
Support 62 → Total 62 NO records

Class 1 (YES / Loan Approved)

Precision 62% → When model says YES, it is correct 62 times out of 100
Recall 65% → Found 65% of actual YES cases
F1-score 64% → Balanced score for YES class
Support 63 → Total 63 YES records

Accuracy = 62% → Model predicted 62 correct out of 100
Macro Avg → Simple average of both classes
Weighted Avg → Average based on number of records

Decision Threshold Analysis

Decision Threshold Analysis means deciding at what probability value the model should say YES or NO.

≥ 0.5 → YES
< 0.5 → NO

```
threshold = 0.5; # Equal importance to YES and NO | Balanced datasets
print("\nDefault Threshold =", threshold)
print("\n Prediction Probabilities (first 10):")
print(y_prob[:10]); # Shows the first 10 probability values predicted by the model
```

Default Threshold = 0.5

```
Prediction Probabilities (first 10):
[0.32716288 0.68291553 0.52806798 0.675654 0.38418969 0.71675992
 0.69488846 0.63420475 0.59271349 0.5549581 ]
```

AUC-ROC Curve

AUC-ROC tells us how well a model can separate positive and negative cases.

It measures how good the model is at distinguishing YES vs NO.

AUC tells how well the model separates Yes vs No.

Step 1: Calculate False Positive Rate, True Positive Rate and Thresholds

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

fpr = how many wrong positives at each threshold

tpr = how many correct positives at each threshold

thresholds = different probability cut-off values

Step 2: Calculate AUC score

auc = roc_auc_score(y_test, y_prob)

AUC tells how well the model separates class 0 and class 1

Higher AUC = better model

```
print("\nAUC Score:", auc)
```

AUC Score: 0.7073732718894009


AUC Score = 0.71

Means 71% chance the model will correctly separate YES vs NO

```
plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, label=f"AUC = {auc:.2f}")
# Plot the ROC curve

plt.plot([0,1], [0,1], '--') # Draws a random guessing line | Used for comparison

plt.xlabel("False Positive Rate");
plt.ylabel("True Positive Rate");
plt.title("ROC Curve")
plt.legend()
plt.grid(True)
plt.show()
```



PREDICTION FUNCTION

```

def predict_loan(age, income, credit, employment):
    employment = le.transform([employment])[0]
    # convert employment text (Salaried/Unemployed/Self-employed)
    # into number

    row = pd.DataFrame([[age, income, credit, employment]],
                       columns=["Age", "Income", "CreditScore",
                           "EmploymentEncoded"])
    # create one-row table exactly like training data

    row = scaler.transform(row)
    # scale values using same scaler used during training

    prob = model.predict_proba(row)[0][1]
    # get probability of loan approval (class = 1)

    print("Probability:", round(prob, 2))
    # print approval probability rounded to 2 decimals

    if prob >= 0.5:
        # check if probability is at least 50%
        print("Loan Approved")
        # approve loan

```

```
else:  
    print("Loan Not Approved")  
    # reject loan  
  
predict_loan(35, 190000, 800, "Salaried");  
print("\n");  
predict_loan(55, 30000, 500, "Unemployed")  
  
Probability: 0.88  
Loan Approved
```

```
Probability: 0.12  
Loan Not Approved
```